


# Shared Memory

# Shared Memory

- Multiple processes accessing the same memory zone
- The memory is accessed as part of the memory zone data of the process
- Advantage is speed; access is instant (no message transfer)



# Shared Memory

- Operations used in shared memory are:
- Creating a Zone.
- Attaching the process area.
- Access to data.
- Release of the area

# Shared Memory

- Functions used in shared memory are:
- `ftok`: identifier(key) generation function.
- `shmget`: creation function.
- `shmat`: attachment process in the area function.
- `shmctl`: control of the area (used for release).

# Shared Memory: The ftok() function

- The ftok() function in C is used to generate a unique key based on a given file path and a project identifier.
- This key is commonly used with other IPC (Inter-Process Communication) mechanisms
  - such as message queues, semaphores, and shared memory segments
- to identify and access resources shared between processes.

Listing 5.11: ftok

```
1 key_t ftok (const char * name, int proj_id);
```

# Shared Memory: The ftok() function

```
include <stdio.h>
#include <sys/types.h>
#include <sys/ipc.h>

int main() {
    char *pathname = "/tmp/example.txt"; // Pathname
    of the file
    int proj_id = 1; // Project identifier

    // Generate a key using ftok()
    key_t key = ftok(pathname, proj_id);
    if (key == -1) {
        perror("ftok");
        return 1;
    }

    printf("Generated key: %d\n", key);

    return 0;
}
```

In this example:

1- We define a pathname variable representing the path to a file.

This file will be used to generate the key.

2- proj\_id is an arbitrary integer identifier for the project.

3- We call ftok() with the pathname and proj\_id as arguments.

4- It returns a key of type key\_t.

The generated key is used to identify a resource shared between processes, such as a message queue, semaphore set, or shared memory segment.

Each resource created with IPC mechanisms like msgget(), semget(), or shmget() is associated with a unique key, allowing processes to access the same resource using that key.

# Shared Memory: shmget() function

- The shmget() function in C is used to create or access a shared memory segment.
- Shared memory is a technique used for inter-process communication, where multiple processes can share a common region of memory.
- This shared memory segment can be used to exchange data between processes efficiently.

Listing 5.12: shmget

```
1 int shmget (key_t key, int size, int shmflg);
```

shmflg takes the values:

- **IPC\_CREAT** which has the effect of creating the zone if it is no exist
- **IPC\_EXCL** returns an error if the zone already exists

# Shared Memory: shmat() function

- The shmat() function in C is used to attach a shared memory segment to the process's address space.
- Once a shared memory segment is attached, the process can access it as if it were a part of its own memory.
- This makes inter-process communication (IPC) via shared memory very efficient as data does not need to be copied between processes.

Listing 5.13: shmat

```
1 void * shmat (int shmid, const void * shmaddr, int shmflg);
```

- **Shmaddr:** put NULL to let the system assign the address.
- **Shmflg:** SHM\_RDONLY for read-only access. Set to 0 for read / write.



# Shared Memory: shmget() function

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>

#define SHM_SIZE 1024 // Size of the shared memory
segment

int main() {
    key_t key = 1234; // Key to identify the shared
memory segment
    int shmid; // Shared memory ID
    char *shmaddr; // Pointer to the shared memory
segment

    // Create or access the shared memory segment
    shmid = shmget(key, SHM_SIZE, IPC_CREAT | 0666);
    if (shmid == -1) {
        perror("shmget");
        exit(1);
    }
}
```

```
// Attach the shared memory segment to the process's
address space
shmaddr = shmat(shmid, NULL, 0);
if (shmaddr == (char *) -1) {
    perror("shmat");
    exit(1);
}

// Write some data to the shared memory
sprintf(shmaddr, "Hello, shared memory!");

// Detach the shared memory segment from the
process's address space
if (shmdt(shmaddr) == -1) {
    perror("shmdt");
    exit(1);
}

return 0;
}
```

# Shared Memory

- `shmaddr = shmat(shmid, NULL, 0);`
- means "attach the shared memory segment identified by `shmid` at an address chosen by the system for reading and writing, and store the address of this attached segment in `shmaddr`."
- Once attached, the processes can use the address pointed to by `shmaddr` to read from and write to the shared memory segment as if they were accessing regular memory.

# Shared Memory: shmctl()

- The shmctl() is used to perform various control operations on a shared memory segment.
- The shmctl() function allows a process to configure and manage shared memory segments based on the command and options provided.
- It's commonly used to change properties of the segment, get information about it, or remove it from the system.

# Shared Memory: shmctl()

shmctl

Listing 5.14: shmctl

```
1 int shmctl (int shmid, int cmd, struct shmid_ds * buf);
```

shmctl returns 0 if Ok.

Values taken by **cmd**:

- **IPC\_SET**: change access permissions (read / write).
- **IPC\_RMID**: suppression of memory.
- **SHM\_LOCK**: not allow switching the memory area.
- **SHM\_UNLOCK**: allow switching.

# Shared Memory: shmget() function

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/shm.h>
#include <sys/ipc.h>

int main() {
    key_t key;
    int shmid;

    // Get key
    key = ftok("path_to_some_file", 'R');

    // Create the shared memory segment
    shmid = shmget(key, 1024, 0644 | IPC_CREAT);
    if (shmid == -1) {
        perror("shmget");
        exit(1);
    }
```

```
// Perform some operation with the shared memory...

// Now remove the shared memory segment
if (shmctl(shmid, IPC_RMID, NULL) == -1) {
    perror("shmctl");
    exit(1);
}

printf("Shared memory segment removed\n");
return 0;
}
```

# Shared Memory: shmctl()

- In this example, the shmctl() function is used with the IPC\_RMID command to mark the shared memory segment for deletion.
- After this call, new processes will not be able to attach to this segment, and it will be destroyed once all attached processes have detached from it.

# Shared Memory: Example

- Develop a C program that simulates the management of a shared resource (e.g., a printer, a computation task scheduler, etc.) among multiple concurrent processes using shared memory.

# Shared Memory: Example

## Specific Tasks

### 1- Initialization Program:

- Use `ftok` to generate a unique key.
- Create a shared memory segment using `shmget`.
- Attach the shared memory segment using `shmat`.
- Initialize the resource status as available in the shared memory.
- Detach the memory segment and leave it available for other processes.

### 2- Client Program:

- Attach to the existing shared memory segment created by the initialization program.
- Check if the resource is available. If not, wait until it becomes available.
- Once available, mark the resource as busy, simulate resource usage (e.g., sleep for a few seconds), and then mark it as free again.
- Finally, detach from the shared memory segment.

### 3- Cleanup Program:

- Attach to the shared memory segment.
- After ensuring no process needs the resource anymore, use `shmctl` with `IPC_RMID` to remove the shared memory segment.



# Shared Memory: Example

1. Initialization Program (init.c): This program creates and initializes the shared memory segment.

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/ipc.h>
#include <sys/shm.h>

#define SHM_KEY_PATH "resource_key"
#define SHM_KEY_ID 65

int main() {
    key_t key;
    int shmid;
    int *resource;

    // Generate unique key
    key = ftok(SHM_KEY_PATH, SHM_KEY_ID);
    if (key == -1) {
        perror("ftok");
        exit(1);
    }
```

```
// Create the shared memory segment
shmid = shmget(key, sizeof(int), 0666 | IPC_CREAT);
if (shmid == -1) {
    perror("shmget");
    exit(1);
}
// Attach the shared memory segment
resource = (int *)shmat(shmid, NULL, 0);
if (resource == (void *)-1) {
    perror("shmat");
    exit(1);
}
// Initialize the resource as available (0)
*resource = 0;

// Detach the shared memory
if (shmdt(resource) == -1) {
    perror("shmdt");
    exit(1);
}
printf("Shared memory segment initialized.\n");
return 0;
}
```

# Shared Memory: Example

2. Client Program (client.c): This program attaches to the shared memory segment, tries to acquire the resource, uses it, and then releases it.

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/ipc.h>
#include <sys/shm.h>

#define SHM_KEY_PATH "resource_key"
#define SHM_KEY_ID 65

int main() {
    key_t key;
    int shmid;
    int *resource;

    // Generate unique key
    key = ftok(SHM_KEY_PATH, SHM_KEY_ID);
    if (key == -1) {
        perror("ftok");
        exit(1);
    }
```

```
// Access the shared memory segment
shmid = shmget(key, sizeof(int), 0666);
if (shmid == -1) {
    perror("shmget");
    exit(1);
}

// Attach the shared memory segment
resource = (int *)shmat(shmid, NULL, 0);
if (resource == (void *)-1) {
    perror("shmat");
    exit(1);
}

// Try to acquire the resource
while (__sync_lock_test_and_set(resource, 1)) {
    printf("Resource busy, waiting...\n");
    sleep(1); // Wait before trying again
}

printf("Resource acquired, using...\n");
sleep(5); // Simulate resource usage
```

# Shared Memory: Example

2. Client Program (client.c): This program attaches to the shared memory segment, tries to acquire the resource, uses it, and then releases it.

```
// Release the resource
__sync_lock_release(resource);
printf("Resource released.\n");

// Detach from the shared memory
if (shmdt(resource) == -1) {
    perror("shmdt");
    exit(1);
}

return 0;
}
```

# Shared Memory: Example

3. Cleanup Program (cleanup.c): This program removes the shared memory segment after all processes are done using it.

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/ipc.h>
#include <sys/shm.h>

#define SHM_KEY_PATH "resource_key"
#define SHM_KEY_ID 65

int main() {
    key_t key;
    int shmid;

    // Generate unique key
    key = ftok(SHM_KEY_PATH, SHM_KEY_ID);
    if (key == -1) {
        perror("ftok");
        exit(1);
    }
}
```

```
// Find the shared memory segment
shmid = shmget(key, sizeof(int), 0666);
if (shmid == -1) {
    perror("shmget");
    exit(1);
}

// Remove the shared memory segment
if (shmctl(shmid, IPC_RMID, NULL) == -1) {
    perror("shmctl");
    exit(1);
}

printf("Shared memory segment removed.\n");
return 0;
}
```

# Shared Memory: Example

- 1- Run `./init` to initialize the shared memory.
- 2- Run multiple `./client` instances concurrently to simulate multiple processes trying to access the resource.
- 3- Once all client processes have finished, run `./cleanup` to clean up the shared memory.